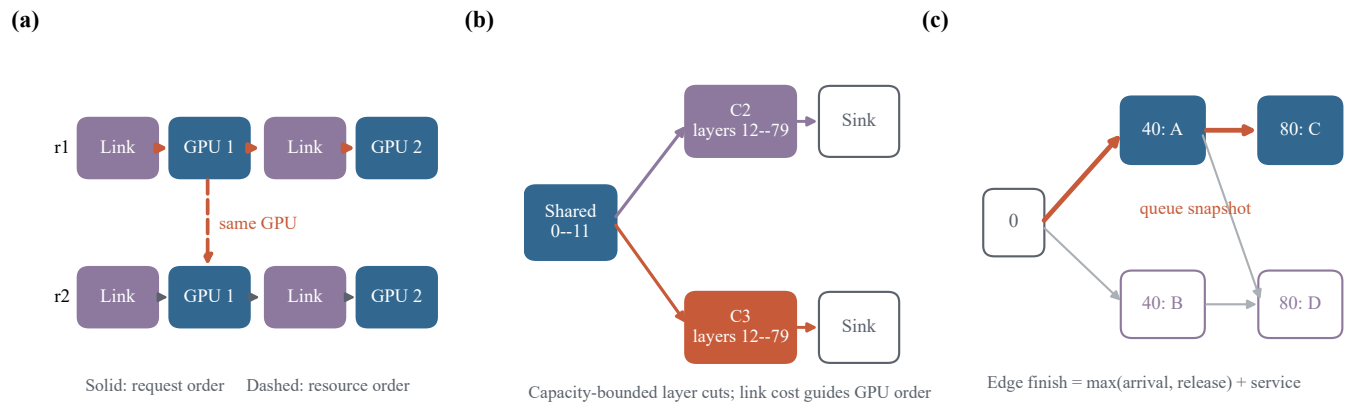
(a)
r1
Link
GPU 1
Link
GPU 2
same GPU
r2
Link
GPU 1
Link
GPU 2
Solid: request order
Dashed: resource order
(b)
Shared
0--11
C2
layers 12--79
Sink
C3
layers 12--79
Sink
Capacity-bounded layer cuts; link cost guides GPU order
(c)
0
40: A
80: C
queue snapshot
40: B
80: D
Edge finish = max(arrival, release) + service