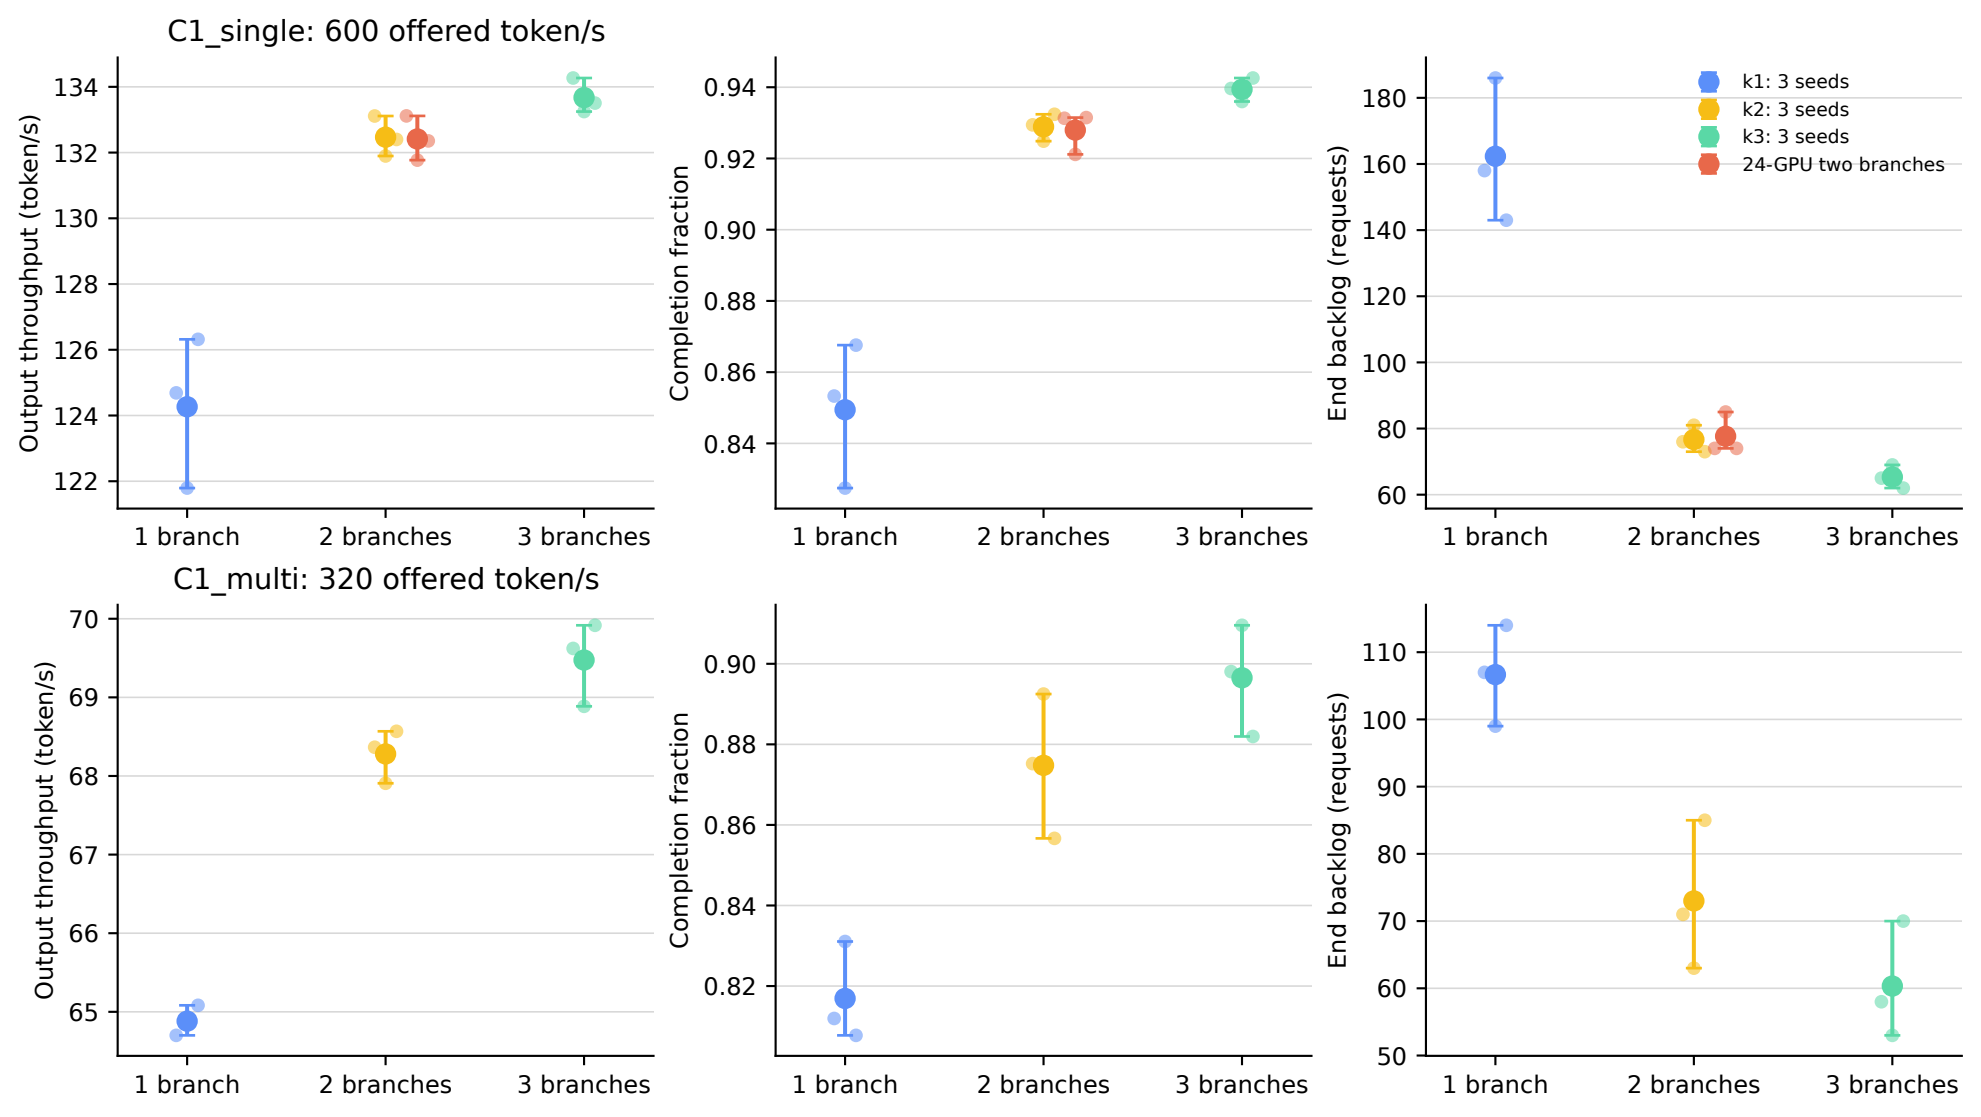
C1_single: 600 offered token/s
Output throughput (token/s)
Completion fraction
End backlog (requests)
k1: 3 seeds
k2: 3 seeds
k3: 3 seeds
24-GPU two branches
1 branch
2 branches
3 branches
C1_multi: 320 offered token/s
Output throughput (token/s)
Completion fraction
End backlog (requests)
1 branch
2 branches
3 branches